

Computer Vision

-Spatial and Frequency Domain Processing-

Oliver Bimber

Course Schedule

Type	Date	Time	Room	Topic	Comment
Lec1	11.10.2022	12:00-13:30	HI	Introduction and Course Overview	
Lab1	10./11./12./13.10.2022	17:15-18:45	S3055	Introduction to Python	
Lec2	18.10.2022	12:00-13:30	HS 1	Spatial and Frequency Domain Processing	
Lab2	17./18./19./20.10.2022	17:15-18:45	S3055	Introduction to IP/CV Modules	
Lec3	25.10.2022	12:00-13:20	HS 1	Gradient Domain Processing	National Holiday (26.10.)
Lec4	08.11.2022	12:00-13:30	HS 10	Segmentation and Local Features	Allerheiligen (2.11.)
Lab3	07./08./09./10.11.2022	17:15-18:45	S3055	Project Introduction	
Lec5	15.11.2022	12:00-13:30	HS 1	Basics of Cameras	
Lec6	22.11.2022	12:00-13:30	HS 1	Geometric Camera Calibration	
Lab4	21./22./23./24.11.2022	17:15-18:45	S3055	Project Basics and Related Work	
Lec7	29.11.2022	12:00-13:30	HS 1	The Geometry of Multiple Views	
Lec8	06.12.2022	12:00-13:30	HS 1	Stereoscopic Depth Estimation	Mariä Empfängnis (8.12.)
Lec9	13.12.2022	12:00-13:30	HS 1	Range Scanning and Data Processing	
Lab5	12./13./14./15.12.2022	17:15-18:45	S3055	Presentation of Initial Ideas	Christmas Break
Lec10	10.01.2023	12:00-13:30	HS 1	Structure from Motion	
Lab6	09./10./11./12.01.2023	17:15-18:45	S3055	Presentation of Intermediate Results and Final Concepts	
Lec11	17.01.2023	12:00-13:30	HS 1	Computational Imaging	
Lec12	24.01.2023	12:00-13:30	HS 1	Recap and Q&A	
Lab7	23./24./25./26.01.2023	17:15-18:45	S3055	Final Project Presentations	
Ex1	31.01.2023	12:00-13:30	HS 1	Exam (Hauptklausur)	
Ex2	28.02.2023	15:30-17:00	TBA	Retry Exam (Nachklausur)	

Image Analysis

- One goal of image analysis is to identify distinct image features
- What are good features?



Image Analysis

- One goal of image analysis is to identify distinct image features
- What are good features?
 - for example, those image points that describe main characteristics of image content, such as edges or corners





Image Analysis


- One goal of image analysis is to identify distinct image features
- What are good features?
 - for example, those image points that describe main characteristics of image content, such as edges or corners
 - but features should also be robust (invariant to rotation, translation, scaling, lighting, etc.), since we want to find the same features if the same content is captured under different conditions
- How do we compute (simple) features?



Linear Filtering

- Which kind of filtering operation is this?

$$R_{i,j} = \frac{1}{(2k+1)^2} \sum_{u=i-k}^{u=i+k} \sum_{v=j-k}^{v=j+k} F_{u,v}$$




linear filtering

Linear Filtering

- Which kind of filtering operation is this?
 - local average
- Image filtering:
 - construct new image R of same size as original image F
 - fill each location in R with a weighted sum of the pixel values of corresponding location's neighbourhood in F
 - use the same weights each time
 - different sets of weights represent different filters
 - shift-invariant: filtered result depends on image neighbourhood – but not on position of neighbourhood
 - linear: filtering the sum of two images is equivalent as summing the two filtered images

$$R_{i,j} = \frac{1}{(2k+1)^2} \sum_{u=i-k}^{u=i+k} \sum_{v=j-k}^{v=j+k} F_{u,v}$$



linear filtering

- Shift invariant linear system (SILS)

Convolution

- The pattern of weights is usually referred to as kernel H
- The process of applying H to an image F is known as convolution
- Say: “The convolution of F with H results in R ”
- The kernel H can be defined as a continuous function or a finite matrix of discrete weights
- What is the kernel for our local average example?

$$R_{i,j} = \frac{1}{(2k+1)^2} \sum_{u=i-k}^{u=i+k} \sum_{v=j-k}^{v=j+k} F_{u,v}$$

$$R_{i,j} = \sum_{u=i-k}^{u=i+k} \sum_{v=j-k}^{v=j+k} H_{i-u, j-v} \cdot F_{u,v}$$

$$R = H \otimes F$$

Convolution

- The pattern of weights is usually referred to as kernel H
- The process of applying H to an image F is known as convolution
- Say: “The convolution of F with H results in R ”
- The kernel H can be defined as a continuous function or a finite matrix of discrete weights
- What is the kernel for our local average example?
- How large is it?

$$R_{i,j} = \frac{1}{(2k+1)^2} \sum_{u=i-k}^{u=i+k} \sum_{v=j-k}^{v=j+k} F_{u,v}$$

$$R_{i,j} = \sum_{u=i-k}^{u=i+k} \sum_{v=j-k}^{v=j+k} H_{i-u, j-v} \cdot F_{u,v}$$

$$R = H \otimes F$$

$$H_{i,j} = \frac{1}{(2k+1)^2}$$

Convolution

- The pattern of weights is usually referred to as kernel H
- The process of applying H to an image F is known as convolution
- Say: “The convolution of F with H results in R ”
- The kernel H can be defined as a continuous function or a finite matrix of discrete weights
- What is the kernel for our local average example?
- How large is it?
 - $(2k+1)^2$ weights

$$\frac{1}{9}$$

1	1	1
1	1	1
1	1	1

“box filter”


$$R_{i,j} = \frac{1}{(2k+1)^2} \sum_{u=i-k}^{u=i+k} \sum_{v=j-k}^{v=j+k} F_{u,v}$$

$$R_{i,j} = \sum_{u=i-k}^{u=i+k} \sum_{v=j-k}^{v=j+k} H_{i-u, j-v} \cdot F_{u,v}$$

$$R = H \otimes F$$

$$H_{i,j} = \frac{1}{(2k+1)^2}$$

Example



Original

0	0	0
0	1	0
0	0	0

?


Example: Identity



Original


0	0	0
0	1	0
0	0	0

?



Original

0	0	0
0	1	0
0	0	0



Filtered
(no change)

Example




Original

 $\frac{1}{9}$

1	1	1
1	1	1
1	1	1

?

Example: Local Average




Original

 $\frac{1}{9}$

1	1	1
1	1	1
1	1	1


?



Original


 $\frac{1}{9}$

1	1	1
1	1	1
1	1	1



Blur (with a
box filter)

Example



Original

0	0	0
0	2	0
0	0	0

-


$\frac{1}{9}$

1	1	1
1	1	1
1	1	1

?

(Note that filter sums to 1)

Example: Sharpening



Original

0	0	0
0	2	0
0	0	0

-

$\frac{1}{9}$

1	1	1
1	1	1
1	1	1

?

(Note that filter sums to 1)



Original

0	0	0
0	2	0
0	0	0

-

$\frac{1}{9}$


1	1	1
1	1	1
1	1	1




Sharpening filter

- Accentuates differences
with local average

Example: Sharpening



before



after

Example: Classification



* ? =



Example: Classification



*




=



Example: Smoothing

- Convoluting an average kernel with an image is known as smoothing or blurring
- Our local average kernel is a smoothing operator
- But it is quite unrealistic (its output does not look like the one of a defocused camera)
- Why?

$$R = H \otimes F$$




$$H_{i,j} = \frac{1}{(2k+1)^2}$$

Example: Smoothing

- Convolving an average kernel with an image is known as smoothing or blurring
- Our local average kernel is a smoothing operator
- But it is quite unrealistic (its output does not look like the one of a defocused camera)
- Why?
 - the intensity spread (point spread function - PSF) of a defocused point is not spatially constant, but follows a Gaussian distribution
- The standard deviation sigma σ controls the influence of the neighbours (small σ - neighbours have small weights)

$$R = H \otimes F$$



$$H_{i,j} = \frac{1}{(2k+1)^2}$$




$$G_{\sigma}(x,y) = \frac{1}{2\pi\sigma^2} \exp\left(-\frac{(x^2 + y^2)}{2\sigma^2}\right)$$

normalization
(sum of weights equals 1)



$$H_{i,j} = \frac{1}{2\pi\sigma^2} \exp\left(-\frac{((i-k-1)^2 + (j-k-1)^2)}{2\sigma^2}\right)$$

Example: Smoothing





Example: Smoothing

- Convolving an average kernel with an image is known as smoothing or blurring
- Our local average kernel is a smoothing operator
- But it is quite unrealistic (its output does not look like the one of a defocused camera)
- Why?
 - the intensity spread (point spread function - PSF) of a defocused point is not spatially constant, but follows a Gaussian distribution
- The standard deviation sigma σ controls the influence of the neighbours (small σ - neighbours have small weights)

$$H_{i,j} =$$

$$\frac{1}{2\pi\sigma^2} \exp\left(-\frac{((i-k-1)^2 + (j-k-1)^2)}{2\sigma^2}\right)$$



smoothing reduces noise

Properties of Gaussians

- Remove “high-frequency” components from the image (low-pass filter)
- Convolution with itself is another Gaussian
 - one can smoothen with small-width kernel multiple times, and get same result as smoothening with larger-width kernel
 - convolving two times with Gaussian kernel of width σ is same as convolving once with kernel of width $\sigma\sqrt{2}$
- Separable kernel
 - factors into product of two 1D Gaussians

2D convolution
(center location only)

1	2	1	*	2	3	3	= 2 + 6 + 3 = 11
2	4	2	*	3	5	5	= 6 + 20 + 10 = 36
1	2	1	*	4	4	6	= 4 + 8 + 6 = 18
							<u>65</u>

The filter factors
into a product of 1D
filters:

1	2	1	=	1	x	1	2	1
2	4	2	=	2				
1	2	1	=	1				

Perform convolution
along rows:

1	2	1	*	2	3	3	=		11	
			*	3	5	5	=		18	
			*	4	4	6	=		18	

Followed by convolution
along the remaining column:

1			*		11		=			
2			*		18		=		65	
1			*		18		=			

Convolution Rules for SILSs

- Superimposition: the sum of two filtered images is equivalent to the filtered image with the sum of the two kernels (component-wise addition)

$$H_1 \otimes R + H_2 \otimes R = (H_1 + H_2) \otimes R$$

- Scaling: scaling of a filtered image is equivalent to filtering the image with a scaled kernel

$$(kH) \otimes R = k(H \otimes R)$$

- Symmetric: the convolution of two filters to an image is symmetric

$$H_1 \otimes (H_2 \otimes R) = H_2 \otimes (H_1 \otimes R)$$

- Associative: convolution is associative, which means that we can find a single kernel that behaves like the composition of two kernels

$$H_1 \otimes (H_2 \otimes R) = (H_1 \otimes H_2) \otimes R$$

Example: Unsharp Masking


<https://en.wikipedia.org/wiki/Lenna>

[https://de.wikipedia.org/wiki/Lena_\(Testbild\)](https://de.wikipedia.org/wiki/Lena_(Testbild))

What does blurring take away?




Let's add it back:




Example: Unsharp Masking

$$\begin{aligned} F + \alpha(F - F * g) &= \\ F * e + \alpha(F * e - F * g) &= \\ F * e + \alpha F * e - \alpha F * g &= \\ (1 + \alpha)(F * e) - \alpha F * g &= \\ F * ((1 + \alpha)e - \alpha g) \end{aligned}$$


0	0	0
0	1	0
0	0	0




unit impulse



Gaussian



Laplacian of Gaussian



Example: Partial Derivations

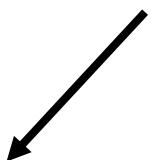
- The partial derivation of a function f can be estimated as a symmetric finite difference
- How does the corresponding filter kernel look like?

$$\frac{\partial f(x,y)}{\partial x} \approx f(x+1,y) - f(x-1,y)$$

Example: Partial Derivations

- The partial derivation of a function f can be estimated as a symmetric finite difference
- How does the corresponding filter kernel look like?
- When is the filter's response large?

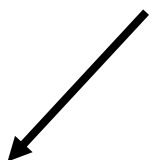
$$\frac{\partial f(x,y)}{\partial x} \approx f(x+1,y) - f(x-1,y)$$


$$H = \begin{bmatrix} 0 & 0 & 0 \\ -1 & 0 & 1 \\ 0 & 0 & 0 \end{bmatrix}$$

Example: Partial Derivations

- The partial derivation of a function f can be estimated as a symmetric finite difference
- How does the corresponding filter kernel look like?
- When is the filter's response large?
 - positive slope (from left to right)
- What about the partial derivation in y direction?

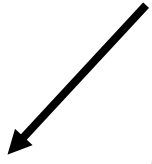
$$\frac{\partial f(x,y)}{\partial x} \approx f(x+1,y) - f(x-1,y)$$


$$H = \begin{bmatrix} 0 & 0 & 0 \\ -1 & 0 & 1 \\ 0 & 0 & 0 \end{bmatrix}$$

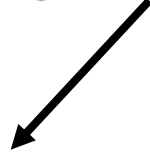
Example: Partial Derivations

- The partial derivation of a function f can be estimated as a symmetric finite difference
- How does the corresponding filter kernel look like?
- When is the filter's response large?
 - positive slope (from left to right)
- What about the partial derivation in y direction?
- Note that finite differences respond strongly to noise!

$$\frac{\partial f(x,y)}{\partial x} \approx f(x+1,y) - f(x-1,y)$$

$$H = \begin{bmatrix} 0 & 0 & 0 \\ -1 & 0 & 1 \\ 0 & 0 & 0 \end{bmatrix}$$


$$\frac{\partial f(x,y)}{\partial y} \approx f(x,y+1) - f(x,y-1)$$

$$H = \begin{bmatrix} 0 & -1 & 0 \\ 0 & 0 & 0 \\ 0 & 1 & 0 \end{bmatrix} \quad H = \begin{bmatrix} 0 & 1 & 0 \\ 0 & 0 & 0 \\ 0 & -1 & 0 \end{bmatrix}$$


Example: Partial Derivations

- The partial derivation of a function f can be estimated as a symmetric finite difference
- How does the corresponding filter kernel look like?
- When is the filter's response large?
 - positive slope (from left to right)
- What about the partial derivation in y direction?
- Note that finite differences respond strongly to noise!



examples for partial derivations in horizontal and vertical directions

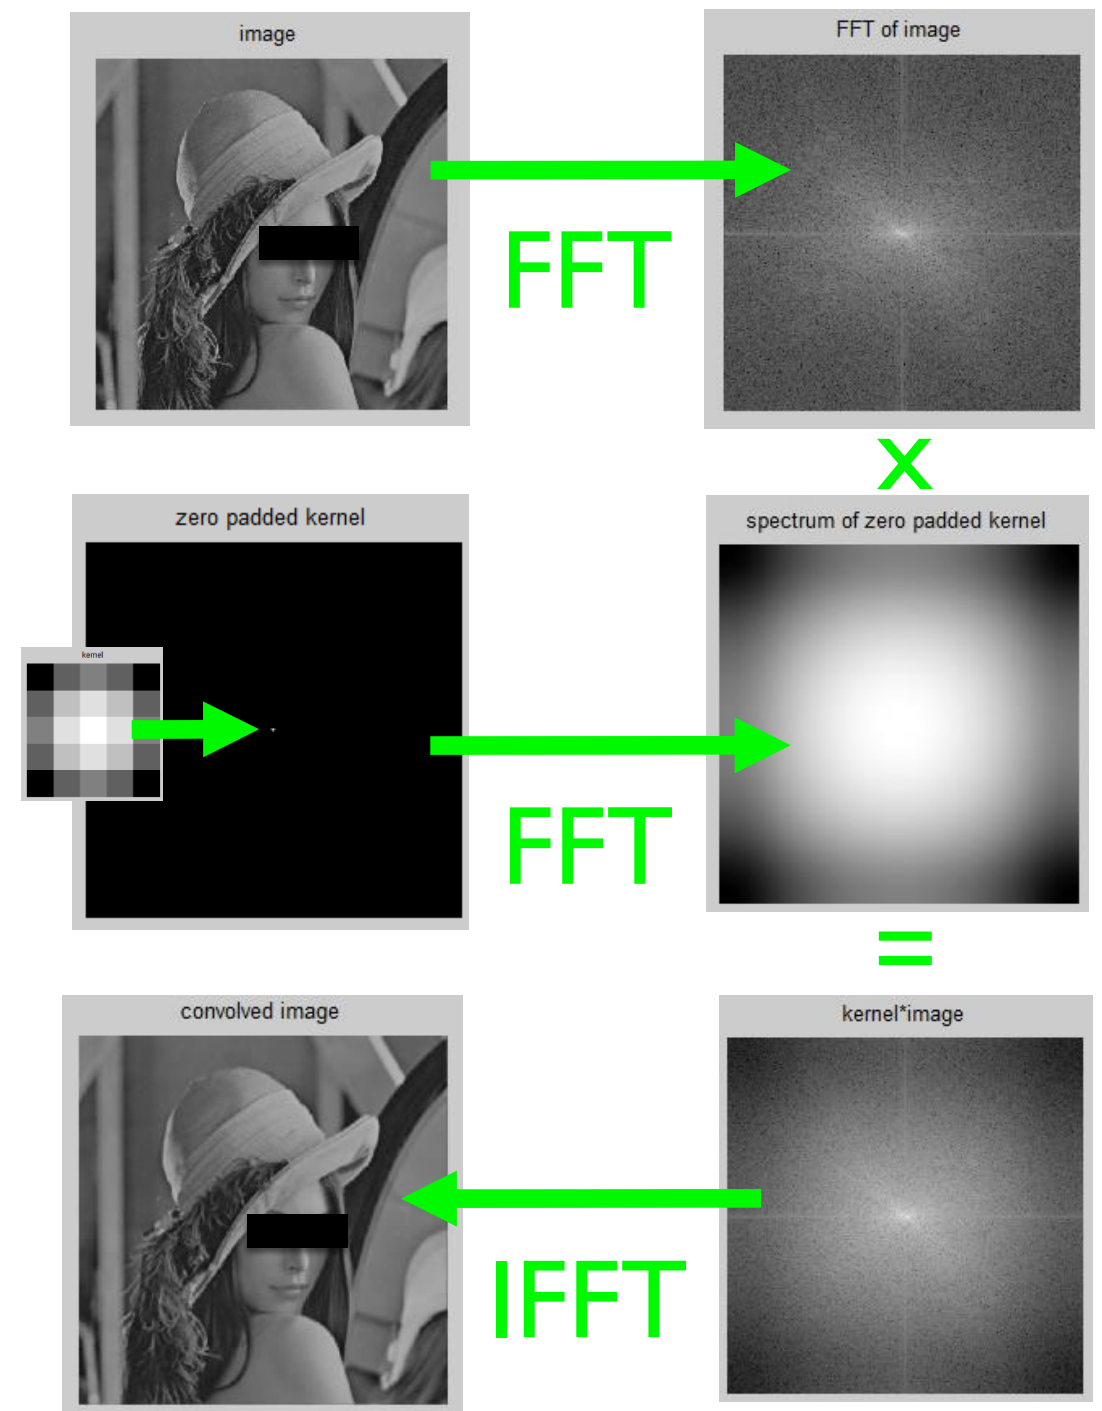
Frequency Decomposition

- An image can also be seen as a discrete two-dimensional signal
- As any signal, it can be decomposed into an integral of weighted basis functions
- A Discrete (Fast) Fourier Transformation applies \cos +/- \sin basis functions at different frequencies, while the weight coefficients represent the amplitude / magnitude of a particular frequency
- A Discrete Cosine Transformation (this is what JPEG compression uses) applies only a cosine basis functions for decomposition
- Thus, every image can be decomposed into a spectrum of frequency of a particular basis function
- This spectrum can be processed / analyzed as well
- The image can be composed from the spectrum through integration




Spatial vs. Frequency Domain

- Convolution theorem:
 - convolution in spatial domain equals a multiplication in frequency domain
 - a division in frequency domain equals the inverse operation, which is called deconvolution / inverse filtering
- In case of deconvolution with a Gaussian kernel, this is equivalent with sharpening (or deblurring)



Deconvolution / Inverse Filtering

- Given an image I that is convolved with signal K results in $I' = I * K + N$ (whereby N is noise)
- If noise N and convolution signal H are known, then original image F can be recovered with $I = K^{-1} * (I' - N)$
- This is known as deconvolution
- The challenge is to estimate K and N
- Example: image deblurring
 - we know that for camera lenses with round apertures the point spread function (PSF) of defocus is Gaussian – thus K is Gaussian
 - the scale of K depends on the amount of defocus
 - noise model can be given or not (then it has to be predicted)
- Wiener filter, regularized filter, Lucy-Richardson algorithm, blind deconvolution



Deconvolution / Inverse Filtering

- Given an image I that is convolved with signal K results in $I' = I * K + N$ (whereby N is noise)
- If noise N and convolution signal H are known, then original image F can be recovered with $I = K^{-1} * (I' - N)$
- This is known as deconvolution
- The challenge is to estimate K and N
- Example: image deblurring
 - we know that for camera lenses with round apertures the point spread function (PSF) of defocus is Gaussian – thus K is Gaussian
 - the scale of K depends on the amount of defocus
 - noise model can be given or not (then it has to be predicted)
- Wiener filter, regularized filter, Lucy-Richardson algorithm, blind deconvolution

convolution in spatial domain:

$$I(x, y) * K_s(x, y) = I'(x, y)$$

convolution in frequency domain

(Fourier transform + convolution theorem):

$$\hat{I}(f_x, f_y) \cdot \hat{K}_s(f_x, f_y) = \hat{I}'(f_x, f_y)$$

deconvolution in frequency domain:

$$\hat{I}(f_x, f_y) = \frac{\hat{I}'(f_x, f_y)}{\hat{K}_s(f_x, f_y)}$$


This is simplified and does not consider a noise model. Division by small values in frequency domain leads to ringing artefacts in spatial domain. Better: apply regularized techniques.



Spatial and Frequency Domain Processing

Filter Banks

- What filters should be used for identifying important features?
 - a lot of research has (and is still) being carried out on this
 - some use filters that are adapted to the human visual system (spots and bars at different orientations)
 - others use simpler edge and line filters
- How do we filter features (e.g., spots or bars) at different scales?




filter bank of two spot and six bar filters

Image Pyramids

- An image pyramid is a collection of representations of the same image at different scales
- Typically, each layer is half the width and half the height of the previous layer
- At each layer, a Gaussian pyramid stores a smoothed (Gaussian kernel) and down sampled version of the previous layer
- A Gauss filter is a low-pass filter
- In terms of signal processing, each level of the Gaussian pyramid is reduced by a subband of higher frequencies (going bottom up)
- The top entry (one pixel) represents the basis (average image intensity)

$$P_G^{n+1}(I) = \downarrow (G_s \otimes P_G^n(I))$$


$$P_G^1(I) = I$$



Gaussian pyramid

Oriented Image Pyramids

- Apply filter bank to levels of pyramid (Gaussian or Laplacian)
- This is called oriented pyramid or steerable pyramid
- It contains the feature response at different scales for Gaussian (or different sub-bands for Laplacian)
- Multiple image pyramids can be the result of a texture analysis process



oriented filter bank of four orientation filters applied to levels of image pyramid pyramid

Decomposition using a Gaussian Basis

- A Gauss filter is a low-pass filter
- Consequently, Gaussian pyramid successively removes high frequencies
- What happens if two consecutive layers are subtracted?

$$P_L^n(I) = P_G^n(I) - \uparrow P_G^{n+1}(I)$$

Decomposition using a Gaussian Basis


- A Gauss filter is a low-pass filter
- Consequently, Gaussian pyramid successively removes high frequencies
- What happens if two consecutive layers are subtracted?
 - the resulting image contains the removed (via Gauss filter) spatial frequencies (a sub-band)
- What happens if this is done for each layer of the Gaussian pyramid?

$$P_L^n(I) = P_G^n(I) - \uparrow P_G^{n+1}(I)$$

Decomposition using a Gaussian Basis

- A Gauss filter is a low-pass filter
- Consequently, Gaussian pyramid successively removes high frequencies
- What happens if two consecutive layers are subtracted?
 - the resulting image contains the removed (via Gauss filter) spatial frequencies (a sub-band)
- What happens if this is done for each layer of the Gaussian pyramid?
 - the result is another image pyramid, called Laplacian pyramid, that can be thought of as the response of a band-pass filter (a sub-band at each level)
- What happens when the pyramid is collapsed?

$$P_L^n(I) = P_G^n(I) - \uparrow P_G^{n+1}(I)$$




Laplacian pyramid

$$P_L^n(I) = P_L^n(I) + \uparrow P_L^{n+1}(I) \rightarrow I$$

Decomposition using a Gaussian Basis

- A Gauss filter is a low-pass filter
- Consequently, Gaussian pyramid successively removes high frequencies
- What happens if two consecutive layers are subtracted?
 - the resulting image contains the removed (via Gauss filter) spatial frequencies (a sub-band)
- What happens if this is done for each layer of the Gaussian pyramid?
 - the result is another image pyramid, called Laplacian pyramid, that can be thought of as the response of a band-pass filter (a sub-band at each level)
- What happens when the pyramid is collapsed?
 - the original image results

$$P_L^n(I) = P_G^n(I) - \uparrow P_G^{n+1}(I)$$



Laplacian pyramid

$$P_L^n(I) = P_L^n(I) + \uparrow P_L^{n+1}(I) \rightarrow I$$

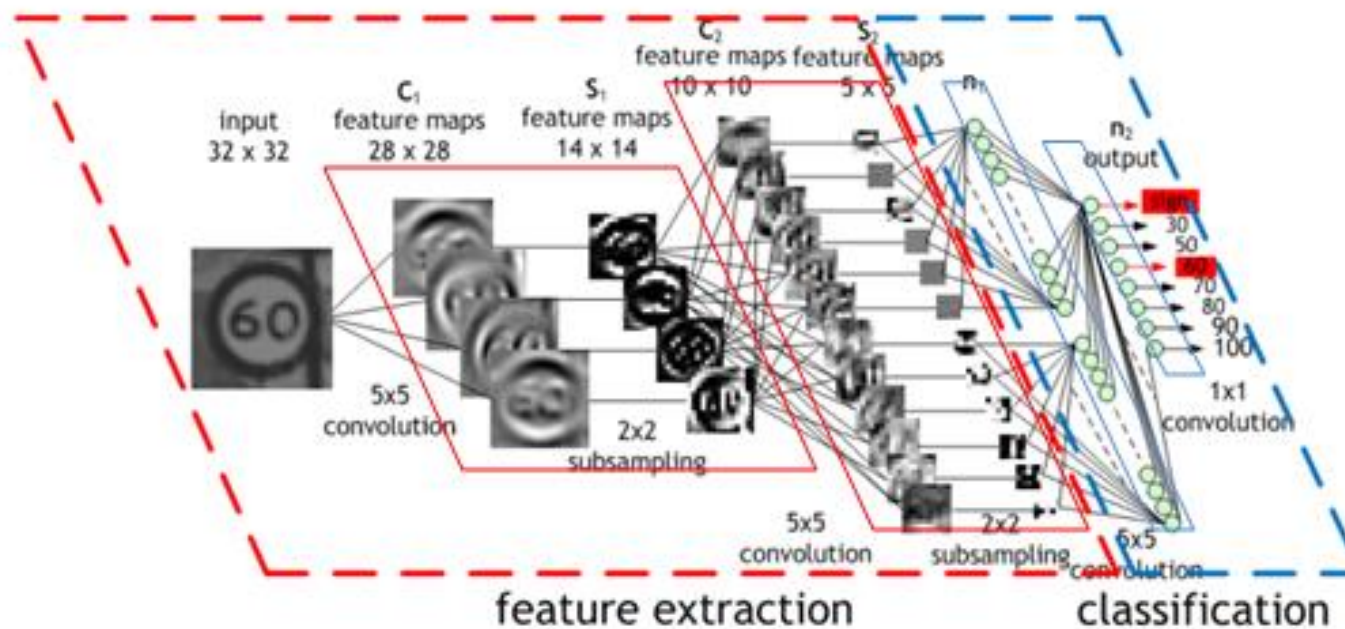
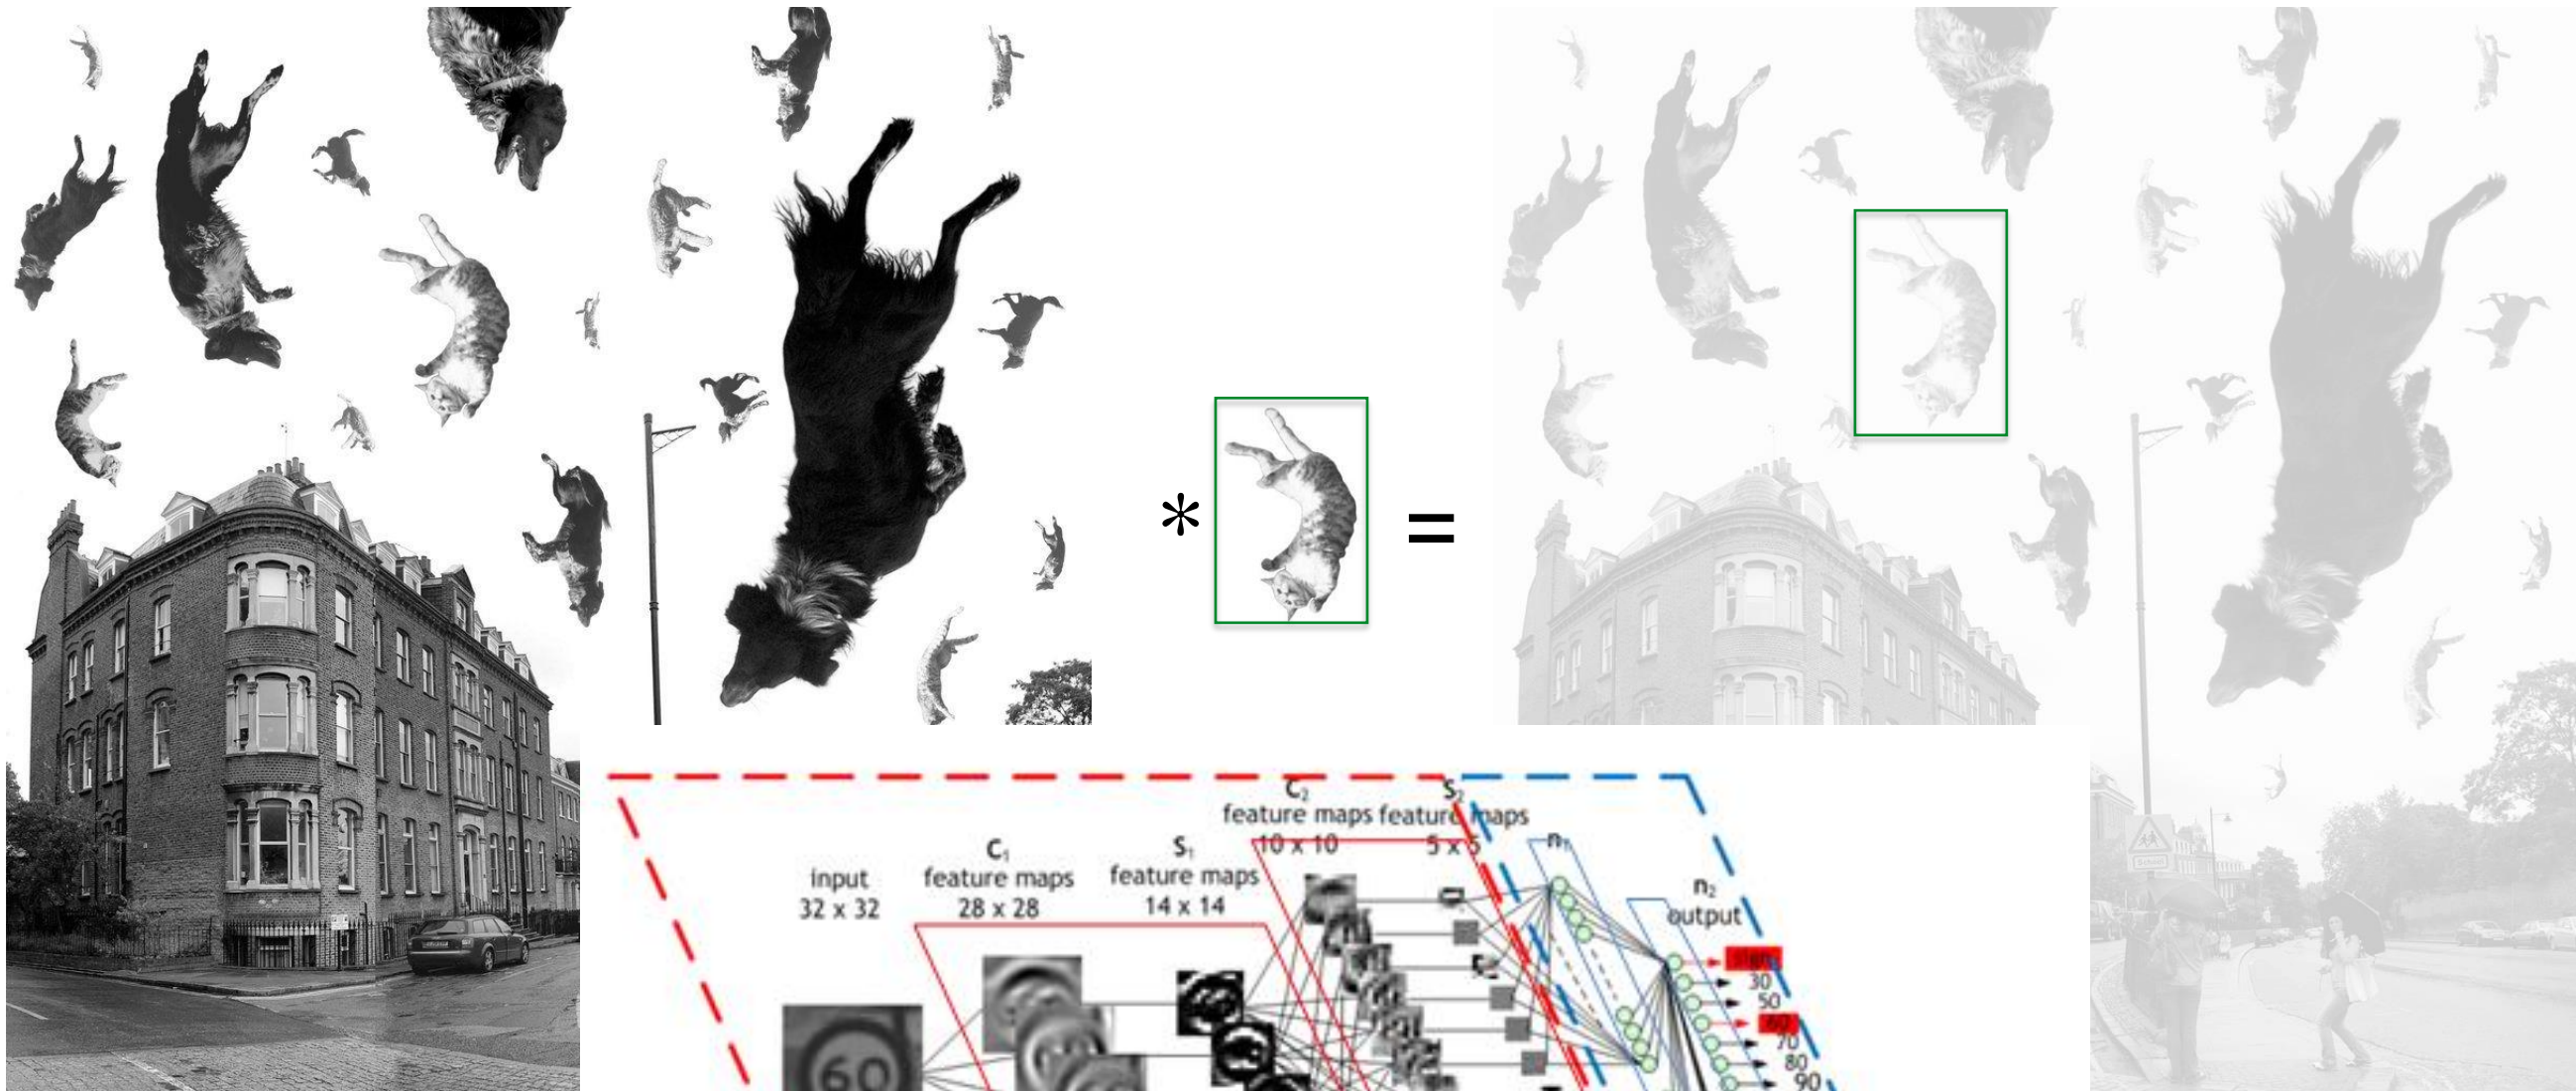
Recap. Example: Classification



$$* \boxed{?} =$$







Recap. Example: Classification








Processing Depth Data

- In many cases, depth maps can be processed like images
- Image operators, such as filters, can be applied
- This, for instance, allows to:
 - handle discontinuities and overlaps
 - fill missing portions
 - smoothen geometric noise
 - analyze depth images
 - find features such as edges
 - segmentation of objects
 - 3D object recognition
 - etc.



Example: LIDAR



Course Schedule

Type	Date	Time	Room	Topic	Comment
Lec1	11.10.2022	12:00-13:30	H1	Introduction and Course Overview	
Lab1	10./11./12./13.10.2022	17:15-18:45	S3055	Introduction to Python	
Lec2	18.10.2022	12:00-13:30	HS 1	Spatial and Frequency Domain Processing	
Lab2	17./18./19./20.10.2022	17:15-18:45	S3055	Introduction to IP/CV Modules	
Lec3	25.10.2022	12:00-13:20	HS 1	Gradient Domain Processing	National Holiday (26.10.)
Lec4	08.11.2022	12:00-13:30	HS 10	Segmentation and Local Features	Allerheiligen (2.11.)
Lab3	07./08./09./10.11.2022	17:15-18:45	S3055	Project Introduction	
Lec5	15.11.2022	12:00-13:30	HS 1	Basics of Cameras	
Lec6	22.11.2022	12:00-13:30	HS 1	Geometric Camera Calibration	
Lab4	21./22./23./24.11.2022	17:15-18:45	S3055	Project Basics and Related Work	
Lec7	29.11.2022	12:00-13:30	HS 1	The Geometry of Multiple Views	
Lec8	06.12.2022	12:00-13:30	HS 1	Stereoscopic Depth Estimation	Mariä Empfängnis (8.12.)
Lec9	13.12.2022	12:00-13:30	HS 1	Range Scanning and Data Processing	
Lab5	12./13./14./15.12.2022	17:15-18:45	S3055	Presentation of Initial Ideas	Christmas Break
Lec10	10.01.2023	12:00-13:30	HS 1	Structure from Motion	
Lab6	09./10./11./12.01.2023	17:15-18:45	S3055	Presentation of Intermediate Results and Final Concepts	
Lec11	17.01.2023	12:00-13:30	HS 1	Computational Imaging	
Lec12	24.01.2023	12:00-13:30	HS 1	Recap and Q&A	
Lab7	23./24./25./26.01.2023	17:15-18:45	S3055	Final Project Presentations	
Ex1	31.01.2023	12:00-13:30	HS 1	Exam (Hauptklausur)	
Ex2	28.02.2023	15:30-17:00	TBA	Retry Exam (Nachklausur)	

Thank You!